

Contents

Declaration of Authorship	i
Abstract	ii
Acknowledgements	iii
Contents	iv
Abbreviations	vii
1 Introduction	1
1.1 How do Artifical Neural Networks compare to the human brain?	1
1.2 Motivation	1
2 Background	2
2.1 Introduction	2
2.2 What are Artifical Neural Networks good at?	2
2.2.1 Classification	3
2.2.2 Generation	3
2.2.3 Reinforcement Learning	3
2.3 What are Artificial Neural Networks bad at?	3
2.3.1 data inefficiency	3
2.4 How to Get Off the Symbol Grounding Merry-Go-Round	3
2.4.1 What is symbol grounding?	3
2.4.2 How do humans do it?	3
2.4.3 How do machines do it?	4
2.5 Why brains are better	4
2.5.1 Embodiment	4
2.5.1.1 Sensory Redundancy	4
2.5.1.2 Biological Filters	4
2.5.2 Development	4
2.5.2.1 Biological Filters, again	4
2.5.3 Pulling yourself up by the bootstraps	4
2.5.4 Machine Equivelancy	4
2.5.4.1 How do we simulate Embodiment for ANNs?	4
2.5.4.2 How do we simulate Development for ANNs?	4

2.6	Where do we go from here?	4
2.6.1	Robot bodies	4
2.6.2	multimodality	4
2.6.3	transfer learning	4
3	A Primer on Artifical Neural Networks	5
3.1	Perceptrons	6
3.1.1	Multi Layer Perceptron	6
3.1.2	Activation Functions	6
3.2	Convolutional Neural Networks	6
3.2.1	What is Convolution?	6
3.2.1.1	Kernels	6
3.2.1.2	Strides	6
3.2.1.3	Dilations	6
3.2.2	Transposed Convolutions	6
3.3	Recurrent Neural Networks	6
3.3.1	Vanilla RNN	6
3.3.2	Gated RNN	6
3.4	Learning Algorithms	6
3.4.1	Gradient Decent	6
3.4.2	Backpropagation	6
3.4.3	Extensions and Improvements	6
3.4.3.1	Learning Rate Schedule	6
3.4.3.2	Momentum	6
3.5	Summary	6
4	Sensory Redundancy: Are two heads better than one?	7
4.1	Why have one sensor when two are better?	7
4.1.1	Aims	7
4.2	Flickr25K	7
4.2.1	dataset description	8
4.2.2	problem description	8
4.2.3	network description	8
4.2.4	results	8
4.2.5	discussion	8
4.3	There and back again: Generating images from natural language	8
4.4	UCU Arabic Spoken Digits and MNIST	8
4.4.1	Dataset Description	8
4.4.1.1	UCU Arabic Spoken Digits	8
4.4.1.2	MNIST Handwritten Digits	9
4.4.2	Problem Description	10
4.4.2.1	Classification	10
4.4.2.2	Bidirectional Symbol Grounding	10
4.5	Experiment Details	11
4.5.1	dataplumbing	11
4.5.1.1	Combining Datasets	11
4.5.2	Merging Modalities	11

4.5.3	Network Description	12
4.5.4	Results	13
4.5.5	Discussion	13
4.6	Conclusion	13
5	Magical Vectors and Where to Find Them	14
5.1	ANN Latent space: the Final Frontier	14
5.1.1	Vector Arithematic Explained	16
5.2	Artificial Shapes Dataset	16
5.2.1	dataset description	16
5.2.2	problem description	16
5.2.3	network description	16
5.2.4	results	16
5.2.4.1	Image Generation	16
5.2.4.2	Multilabel classification	16
5.2.4.3	Vector Arithematic	16
5.2.5	discussion	16
5.3	Real Shapes Dataset	16
5.3.1	dataset description	16
5.3.2	problem description	16
5.3.3	network description	16
5.3.4	results	16
5.3.4.1	Image Generation	16
5.3.4.2	Multilabel classification	16
5.3.4.3	Vector Arithematic	16
5.3.5	discussion	16
6	Transfer Learning: kick-starting the learning process	17
6.1	What is transfer learning	17
6.1.1	A biological analogue	17
6.2	The importance of transfer learning	17
6.3	MultiSense 1 dataset	17
6.3.1	dataset description	17
6.3.2	problem description: image generation from raw speech audio . . .	17
6.3.3	network description	17
6.3.4	results	17
6.3.5	discussion	17
7	What have we learnt and more importantly, what have our neural nets learnt?	18
7.1	Summary of important points	18
7.2	Conclusion	18
8	Future Work	19
8.1	Where do we go now?	19
8.1.1	More Sensors	19
8.1.2	Better Sensors	19
8.1.3	Swarm Robots and the Cloud	19